

# NFS-like system on FUSE

Adithya Bhat - adbhat@cs.wisc.edu

## Setup

FUSE 2.9.7

Apache Thrift 0.9.3

Ubuntu VMs - 2 vCPUs, 7.5GB memory

## Design

### Layers

FUSE - RPCGateway - Thrift Client ==== Thrift Server - RPCGateway - NFS Server.

Pro - readability, maintainability, easy to extend.

Con - copying between layers, not sure if pass by reference mitigates this problem.

### Connection Persistence

Pro - better performance

Con - harder crash handling, cannot handle multiple connections per thread

In a throughput > latency case, would argue that connections are not optimal.

### Commit

Buffers on the server side

NOTE : I thought it would be more interesting to buffer all the data until and unless an fsync or a release is called, so implemented that. The complication here is discussed under issues.

Other Classes - WriteCache and WriteCacheEntry.

Reason - to abstract out the collection buffering the write requests

Why ? - make the insert, flush operations thread-safe. (Multi-threaded server)

Con - costly in terms of time and space.

Unless the application really knows what it's doing, seems sub-optimal.

### Piece by piece lookup

Unnecessary / Implicitly handled by FUSE ?

## Crash Consistency

Server Crashes -

Retry with exponential backoff, 10 tries, starts at 1 second.

## File Handle

stat->dev\_id, stat->ino

ioctl call to get generation number

Optionally, can use the path also here

NOTE : generation numbers are huge.. Instead of combining by adding, etc, can combine by selecting how many digits we think would be necessary. Very unlikely that one inode will see a million generations over one NFS session (right?)

## Issues

### Commit Semantics

Situation : multiple clients accessing the same file.

Server is stateless, does not differentiate between clients.

Fsync called from any client will cause flushing of buffer

How can Client1 tell the difference between fsync by another client and server crash ?

## Versions

### V2

Not a persistent connection

No explicit write caching

### V3

Added COMMIT enhancement, and multithreaded server

### V4

Persistent connection + retry











